

DLCV HW2 Report

資工所 呂兆凱 R11922098

GAN

1. Please print the model architecture of method A and B.

method A : (DCGAN)

Discriminator

```
Discriminator(  
  (main): Sequential(  
    (0): Conv2d(3, 64, kernel_size=(4, 4), stride=(2, 2), padding=(1, 1), bias=False)  
    (1): LeakyReLU(negative_slope=0.2, inplace=True)  
    (2): Conv2d(64, 128, kernel_size=(4, 4), stride=(2, 2), padding=(1, 1), bias=False)  
    (3): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)  
    (4): LeakyReLU(negative_slope=0.2, inplace=True)  
    (5): Conv2d(128, 256, kernel_size=(4, 4), stride=(2, 2), padding=(1, 1), bias=False)  
    (6): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)  
    (7): LeakyReLU(negative_slope=0.2, inplace=True)  
    (8): Conv2d(256, 512, kernel_size=(4, 4), stride=(2, 2), padding=(1, 1), bias=False)  
    (9): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)  
    (10): LeakyReLU(negative_slope=0.2, inplace=True)  
    (11): Conv2d(512, 1, kernel_size=(4, 4), stride=(1, 1), bias=False)  
    (12): Sigmoid()  
  )  
)
```

Generator

```
Generator(  
  (main): Sequential(  
    (0): ConvTranspose2d(100, 512, kernel_size=(4, 4), stride=(1, 1), bias=False)  
    (1): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)  
    (2): ReLU(inplace=True)  
    (3): ConvTranspose2d(512, 256, kernel_size=(4, 4), stride=(2, 2), padding=(1, 1), bias=False)  
    (4): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)  
    (5): ReLU(inplace=True)  
    (6): ConvTranspose2d(256, 128, kernel_size=(4, 4), stride=(2, 2), padding=(1, 1), bias=False)  
    (7): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)  
    (8): ReLU(inplace=True)  
    (9): ConvTranspose2d(128, 64, kernel_size=(4, 4), stride=(2, 2), padding=(1, 1), bias=False)  
    (10): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)  
    (11): ReLU(inplace=True)  
    (12): ConvTranspose2d(64, 3, kernel_size=(4, 4), stride=(2, 2), padding=(1, 1), bias=False)  
    (13): Tanh()  
  )  
)
```

method B : (SAGAN)

Discriminator

```
Discriminator(  
  (14): Sequential(  
    (0): SpectralNorm(  
      (module): Conv2d(256, 512, kernel_size=(4, 4), stride=(2, 2), padding=(1, 1))  
    )  
    (1): LeakyReLU(negative_slope=0.1)  
  )  
  (11): Sequential(  
    (0): SpectralNorm(  
      (module): Conv2d(3, 64, kernel_size=(4, 4), stride=(2, 2), padding=(1, 1))  
    )  
    (1): LeakyReLU(negative_slope=0.1)  
  )  
  (12): Sequential(  
    (0): SpectralNorm(  
      (module): Conv2d(64, 128, kernel_size=(4, 4), stride=(2, 2), padding=(1, 1))  
    )  
    (1): LeakyReLU(negative_slope=0.1)  
  )  
  (13): Sequential(  
    (0): SpectralNorm(  
      (module): Conv2d(128, 256, kernel_size=(4, 4), stride=(2, 2), padding=(1, 1))  
    )  
    (1): LeakyReLU(negative_slope=0.1)  
  )  
  (last): Sequential(  
    (0): Conv2d(512, 1, kernel_size=(4, 4), stride=(1, 1))  
  )  
  (attn1): Self_Attn(  
    (query_conv): Conv2d(256, 32, kernel_size=(1, 1), stride=(1, 1))  
    (key_conv): Conv2d(256, 32, kernel_size=(1, 1), stride=(1, 1))  
    (value_conv): Conv2d(256, 256, kernel_size=(1, 1), stride=(1, 1))  
    (softmax): Softmax(dim=-1)  
  )  
  (attn2): Self_Attn(  
    (query_conv): Conv2d(512, 64, kernel_size=(1, 1), stride=(1, 1))  
  
    (key_conv): Conv2d(512, 64, kernel_size=(1, 1), stride=(1, 1))  
    (value_conv): Conv2d(512, 512, kernel_size=(1, 1), stride=(1, 1))  
    (softmax): Softmax(dim=-1)  
  )  
)
```

Generator

```
Generator(  
  (14): Sequential(  
    (0): SpectralNorm(  
      (module): ConvTranspose2d(128, 64, kernel_size=(4, 4), stride=(2, 2), padding=(1, 1))  
    )  
    (1): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)  
    (2): ReLU()  
  )  
  (11): Sequential(  
    (0): SpectralNorm(  
      (module): ConvTranspose2d(128, 512, kernel_size=(4, 4), stride=(1, 1))  
    )  
    (1): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)  
    (2): ReLU()  
  )  
  (12): Sequential(  
    (0): SpectralNorm(  
      (module): ConvTranspose2d(512, 256, kernel_size=(4, 4), stride=(2, 2), padding=(1, 1))  
    )  
    (1): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)  
    (2): ReLU()  
  )  
  (13): Sequential(  
    (0): SpectralNorm(  
      (module): ConvTranspose2d(256, 128, kernel_size=(4, 4), stride=(2, 2), padding=(1, 1))  
    )  
    (1): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)  
    (2): ReLU()  
  )  
  (last): Sequential(  
    (0): ConvTranspose2d(64, 3, kernel_size=(4, 4), stride=(2, 2), padding=(1, 1))  
    (1): Tanh()  
  )  
  (attn1): Self_Attn(  
    (query_conv): Conv2d(128, 16, kernel_size=(1, 1), stride=(1, 1))  
    (key_conv): Conv2d(128, 16, kernel_size=(1, 1), stride=(1, 1))  
    (value_conv): Conv2d(128, 128, kernel_size=(1, 1), stride=(1, 1))  
    (softmax): Softmax(dim=-1)  
  )  
  (attn2): Self_Attn(  
    (query_conv): Conv2d(64, 8, kernel_size=(1, 1), stride=(1, 1))  
    (key_conv): Conv2d(64, 8, kernel_size=(1, 1), stride=(1, 1))  
    (value_conv): Conv2d(64, 64, kernel_size=(1, 1), stride=(1, 1))  
    (softmax): Softmax(dim=-1)  
  )  
)
```

2. Please show the first 32 generated images of both method A and B then discuss the difference between method A and B.

method A : (DCGAN)



method B : (SAGAN)



The difference between method A and B :

DCGAN是單純利用兩個CNN來分別當作Generator和Discriminator, 而SAGAN則又再加入了pixelwise的Self-Attention機制, 可以更準確的對全局的圖像做調整, 使圖片可以產生得更為細緻。看圖可以發現 SAGAN 所產出的圖片成果更佳, 且其產生的背景也更為清楚明亮。

3. Please discuss what you've observed and learned from implementing GAN.

GAN 其實不好訓練, 因為我們必須讓 Generator 和 Discriminator 的 loss 一起下降, 讓他們一起變好, 才可以達到最好的效果, 若是有一方過於強勢, 都可能導致無法訓練起來的情況。而我們可以透過調整learning rate或是加入weight_decay 的方式來使兩者共同進步以增進其表現。

Diffusion models

1. Please print your model architecture and describe your implementation details.

model architecture

```
UNet_conditional(
  (inc): DoubleConv(
    (double_conv): Sequential(
      (0): Conv2d(3, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
      (1): GroupNorm(1, 64, eps=1e-05, affine=True)
      (2): GELU(approximate=none)
      (3): Conv2d(64, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
      (4): GroupNorm(1, 64, eps=1e-05, affine=True)
    )
  )
  (down1): Down(
    (maxpool_conv): Sequential(
      (0): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)
      (1): DoubleConv(
        (double_conv): Sequential(
          (0): Conv2d(64, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
          (1): GroupNorm(1, 64, eps=1e-05, affine=True)
          (2): GELU(approximate=none)
          (3): Conv2d(64, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
          (4): GroupNorm(1, 64, eps=1e-05, affine=True)
        )
      )
    )
  )
  (2): DoubleConv(
    (double_conv): Sequential(
      (0): Conv2d(64, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
      (1): GroupNorm(1, 128, eps=1e-05, affine=True)
      (2): GELU(approximate=none)
      (3): Conv2d(128, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
      (4): GroupNorm(1, 128, eps=1e-05, affine=True)
    )
  )
  (emb_layer): Sequential(
    (0): SiLU()
    (1): Linear(in_features=256, out_features=128, bias=True)
  )
  (sa1): SelfAttention(
    (mha): MultiheadAttention(
      (out_proj): NonDynamicallyQuantizableLinear(in_features=128, out_features=128, bias=True)
    )
    (ln): LayerNorm((128,), eps=1e-05, elementwise_affine=True)
    (ff_self): Sequential(
      (0): LayerNorm((128,), eps=1e-05, elementwise_affine=True)
      (1): Linear(in_features=128, out_features=128, bias=True)
      (2): GELU(approximate=none)
      (3): Linear(in_features=128, out_features=128, bias=True)
    )
  )
)
```

```

(down2): Down(
  (maxpool_conv): Sequential(
    (0): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)
    (1): DoubleConv(
      (double_conv): Sequential(
        (0): Conv2d(128, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
        (1): GroupNorm(1, 128, eps=1e-05, affine=True)
        (2): GELU(approximate=none)
        (3): Conv2d(128, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
        (4): GroupNorm(1, 128, eps=1e-05, affine=True)
      )
    )
  )
  (2): DoubleConv(
    (double_conv): Sequential(
      (0): Conv2d(128, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
      (1): GroupNorm(1, 256, eps=1e-05, affine=True)
      (2): GELU(approximate=none)
      (3): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
      (4): GroupNorm(1, 256, eps=1e-05, affine=True)
    )
  )
)
(emb_layer): Sequential(
  (0): SiLU()
  (1): Linear(in_features=256, out_features=256, bias=True)
)
)
(sa2): SelfAttention(
  (mha): MultiheadAttention(
    (out_proj): NonDynamicallyQuantizableLinear(in_features=256, out_features=256, bias=True)
  )
  (ln): LayerNorm((256,), eps=1e-05, elementwise_affine=True)
  (ff_self): Sequential(
    (0): LayerNorm((256,), eps=1e-05, elementwise_affine=True)
    (1): Linear(in_features=256, out_features=256, bias=True)
    (2): GELU(approximate=none)
    (3): Linear(in_features=256, out_features=256, bias=True)
  )
)
)
(down3): Down(
  (maxpool_conv): Sequential(
    (0): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)
    (1): DoubleConv(
      (double_conv): Sequential(
        (0): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
        (1): GroupNorm(1, 256, eps=1e-05, affine=True)
        (2): GELU(approximate=none)
        (3): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
        (4): GroupNorm(1, 256, eps=1e-05, affine=True)
      )
    )
  )
  (2): DoubleConv(
    (double_conv): Sequential(
      (0): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
      (1): GroupNorm(1, 256, eps=1e-05, affine=True)
    )
  )
)

```

```

        (2): GELU(approximate=none)
        (3): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
        (4): GroupNorm(1, 256, eps=1e-05, affine=True)
    )
)
)
(emb_layer): Sequential(
  (0): SiLU()
  (1): Linear(in_features=256, out_features=256, bias=True)
)
)
(sa3): SelfAttention(
  (mha): MultiheadAttention(
    (out_proj): NonDynamicallyQuantizableLinear(in_features=256, out_features=256, bias=True)
  )
  (ln): LayerNorm((256,), eps=1e-05, elementwise_affine=True)
  (ff_self): Sequential(
    (0): LayerNorm((256,), eps=1e-05, elementwise_affine=True)
    (1): Linear(in_features=256, out_features=256, bias=True)
    (2): GELU(approximate=none)
    (3): Linear(in_features=256, out_features=256, bias=True)
  )
)
)
(bot1): DoubleConv(
  (double_conv): Sequential(
    (0): Conv2d(256, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
    (1): GroupNorm(1, 512, eps=1e-05, affine=True)
    (2): GELU(approximate=none)
    (3): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
    (4): GroupNorm(1, 512, eps=1e-05, affine=True)
  )
)
)
(bot2): DoubleConv(
  (double_conv): Sequential(
    (0): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
    (1): GroupNorm(1, 512, eps=1e-05, affine=True)
    (2): GELU(approximate=none)
    (3): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
    (4): GroupNorm(1, 512, eps=1e-05, affine=True)
  )
)
)
(bot3): DoubleConv(
  (double_conv): Sequential(
    (0): Conv2d(512, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
    (1): GroupNorm(1, 256, eps=1e-05, affine=True)
    (2): GELU(approximate=none)
    (3): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
    (4): GroupNorm(1, 256, eps=1e-05, affine=True)
  )
)
)
(up1): Up(
  (up1): Upsample(scale_factor=2.0, mode=bilinear)
  (up2): Upsample(scale_factor=2.5, mode=bilinear)
  (conv): Sequential(
    (0): DoubleConv(

```

```

(double_conv): Sequential(
  (0): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
  (1): GroupNorm(1, 512, eps=1e-05, affine=True)
  (2): GELU(approximate=none)
  (3): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
  (4): GroupNorm(1, 512, eps=1e-05, affine=True)
)
)
(1): DoubleConv(
  (double_conv): Sequential(
    (0): Conv2d(512, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
    (1): GroupNorm(1, 256, eps=1e-05, affine=True)
    (2): GELU(approximate=none)
    (3): Conv2d(256, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
    (4): GroupNorm(1, 128, eps=1e-05, affine=True)
  )
)
)
(emb_layer): Sequential(
  (0): SiLU()
  (1): Linear(in_features=256, out_features=128, bias=True)
)
)
(sa4): SelfAttention(
  (mha): MultiheadAttention(
    (out_proj): NonDynamicallyQuantizableLinear(in_features=128, out_features=128, bias=True)
  )
  (ln): LayerNorm((128,), eps=1e-05, elementwise_affine=True)
  (ff_self): Sequential(
    (0): LayerNorm((128,), eps=1e-05, elementwise_affine=True)
    (1): Linear(in_features=128, out_features=128, bias=True)
    (2): GELU(approximate=none)
    (3): Linear(in_features=128, out_features=128, bias=True)
  )
)
)
(up2): Up(
  (up1): Upsample(scale_factor=2.0, mode=bilinear)
  (up2): Upsample(scale_factor=2.5, mode=bilinear)
  (conv): Sequential(
    (0): DoubleConv(
      (double_conv): Sequential(
        (0): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
        (1): GroupNorm(1, 256, eps=1e-05, affine=True)
        (2): GELU(approximate=none)
        (3): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
        (4): GroupNorm(1, 256, eps=1e-05, affine=True)
      )
    )
  )
  (1): DoubleConv(
    (double_conv): Sequential(
      (0): Conv2d(256, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
      (1): GroupNorm(1, 128, eps=1e-05, affine=True)
      (2): GELU(approximate=none)
      (3): Conv2d(128, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
      (4): GroupNorm(1, 64, eps=1e-05, affine=True)
    )
  )
)

```



```

    )
    )
    )
    (emb_layer): Sequential(
      (0): SiLU()
      (1): Linear(in_features=256, out_features=64, bias=True)
    )
  )
  (sa5): SelfAttention(
    (mha): MultiheadAttention(
      (out_proj): NonDynamicallyQuantizableLinear(in_features=64, out_features=64, bias=True)
    )
    (ln): LayerNorm((64,), eps=1e-05, elementwise_affine=True)
    (ff_self): Sequential(
      (0): LayerNorm((64,), eps=1e-05, elementwise_affine=True)
      (1): Linear(in_features=64, out_features=64, bias=True)
      (2): GELU(approximate=none)
      (3): Linear(in_features=64, out_features=64, bias=True)
    )
  )
  (up3): Up(
    (up1): Upsample(scale_factor=2.0, mode=bilinear)
    (up2): Upsample(scale_factor=2.5, mode=bilinear)
    (conv): Sequential(
      (0): DoubleConv(
        (double_conv): Sequential(
          (0): Conv2d(128, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
          (1): GroupNorm(1, 128, eps=1e-05, affine=True)
          (2): GELU(approximate=none)
          (3): Conv2d(128, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
          (4): GroupNorm(1, 128, eps=1e-05, affine=True)
        )
      )
      (1): DoubleConv(
        (double_conv): Sequential(
          (0): Conv2d(128, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
          (1): GroupNorm(1, 64, eps=1e-05, affine=True)
          (2): GELU(approximate=none)
          (3): Conv2d(64, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
          (4): GroupNorm(1, 64, eps=1e-05, affine=True)
        )
      )
    )
  )
  (emb_layer): Sequential(
    (0): SiLU()
    (1): Linear(in_features=256, out_features=64, bias=True)
  )
  (sa6): SelfAttention(
    (mha): MultiheadAttention(
      (out_proj): NonDynamicallyQuantizableLinear(in_features=64, out_features=64, bias=True)
    )
    (ln): LayerNorm((64,), eps=1e-05, elementwise_affine=True)
    (ff_self): Sequential(
      (0): LayerNorm((64,), eps=1e-05, elementwise_affine=True)

```

)

implementation details

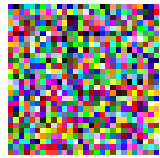
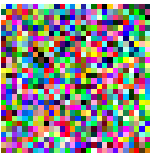
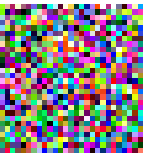
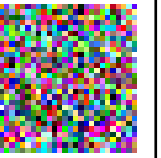
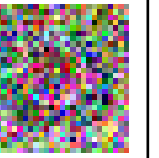

Diffusion 用來對於圖片逐步增加noise, 使其漸進變換為一張 noise 圖片; 並在訓練過程中使用一個 U-Net 結構的 Autoencoder, 來對於在 t 時間的 noise 進行預測並去噪。最後我們就可以使用此U-Net對於隨機亂數去噪產生圖片。

About Parameter : batch_size : 128, learning rate : $3 * 10^{-3}$, 使用 AdamW optimizer

2. Please show 10 generated images for each digit (0-9) in your report. You can put all 100 outputs in one image with columns indicating different noise inputs and rows indicating different digits.



- Visualize total six images in the reverse process of the first “0” in your grid in (2) with different time steps.

					
t = 0	t = 200	t = 400	t = 600	t = 800	t = 1000

- Please discuss what you’ve observed and learned from implementing conditional diffusion model.

認識到除了 GAN 之外的影像生成模型，並且也學到了 diffusion model 的運作原理與實作方式，也覺得的作者想法很有趣，同樣的方法也可以應用在其他領域如 Segmentation、Super Resolution 等。另外也發現雖然 diffusion model 的 sample 速度很慢，但可以生成到品質很好的圖像。

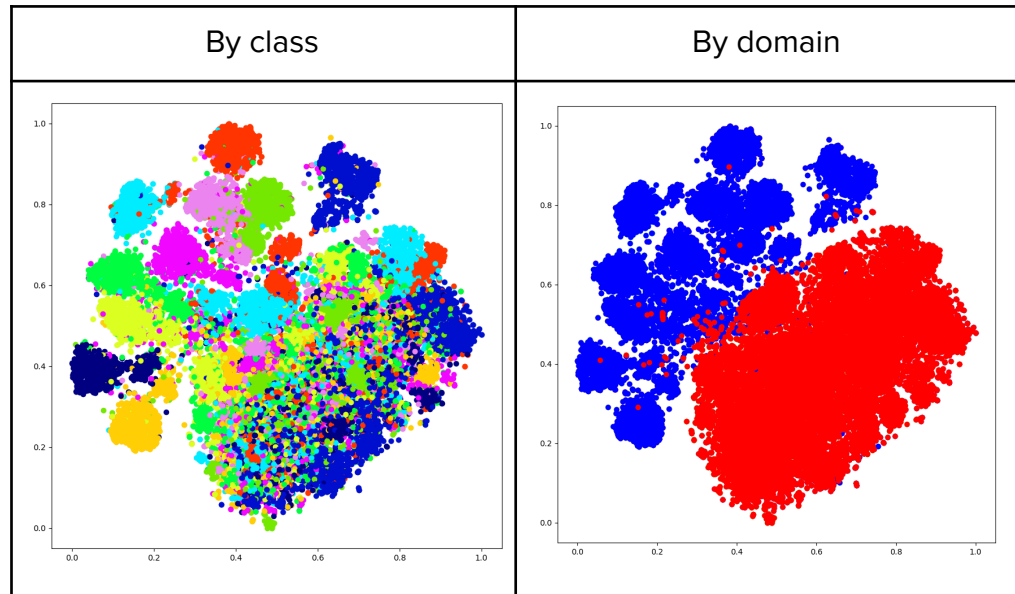
DANN

- Please create and fill the table with the following format in your report:


	MNIST-M → SVHN	MNIST-M → USPS
Trained on source	39%	80%
Adaptation (DANN)	49%	82%
Trained on target	92%	99%

- Please visualize the latent space of DANN by mapping the validation images to 2D space with t-SNE. For each scenario, you need to plot two figures which are colored by digit class (0-9) and by domain, respectively

1. MNIST-M \rightarrow SVHN



2. MNIST-M \rightarrow USPS



3. Please describe the implementation details of your model and discuss what you've observed and learned from implementing DANN.

DANN model :

```
DANN(  
  (feature): Sequential(  
    (0): Conv2d(3, 32, kernel_size=(5, 5), stride=(1, 1))  
    (1): BatchNorm2d(32, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)  
    (2): ReLU(inplace=True)  
    (3): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)  
    (4): Conv2d(32, 48, kernel_size=(5, 5), stride=(1, 1))  
    (5): BatchNorm2d(48, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)  
    (6): Dropout2d(p=0.5, inplace=False)  
    (7): ReLU(inplace=True)  
    (8): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)  
  )  
  (classifier): Sequential(  
    (0): Linear(in_features=768, out_features=384, bias=True)  
    (1): BatchNorm1d(384, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)  
    (2): ReLU(inplace=True)  
    (3): Linear(in_features=384, out_features=192, bias=True)  
    (4): BatchNorm1d(192, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)  
    (5): ReLU(inplace=True)  
    (6): Linear(in_features=192, out_features=10, bias=True)  
  )  
  (discriminator): Sequential(  
    (0): Linear(in_features=768, out_features=384, bias=True)  
    (1): BatchNorm1d(384, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)  
    (2): ReLU(inplace=True)  
    (3): Linear(in_features=384, out_features=2, bias=True)  
  )  
)
```

the implementation details of your model and discuss what you've observed and learned :

3 components : feature extractor, classifier, discriminator

About Parameter : batch_size : 128, learning rate : 10^{-3} , 使用 Adam optimizer

利用feature extractor取出feature, 並將其分別給予classifier和discriminator, 並在discriminator前面加一個GradReverse做梯度反轉。

發現使用 DANN 在 MNIST-M \rightarrow USPS 的結果上比 MNIST-M \rightarrow SVHN 的結果來得好很多, 因為 USPS 為黑白圖片資料集, 而 source domain 為 MNIST-M, MNIST-M 是彩色資料帶有較多資訊, 因此可以讓預測在黑白圖片上的效果較好; 而若是預測在同樣為彩色資料的SVHN, 效果就沒有那麼好。

Reference

- GAN :
 - <https://arxiv.org/abs/1511.06434>
 - https://pytorch.org/tutorials/beginner/dcgan_faces_tutorial.html
 - <https://ithelp.ithome.com.tw/articles/10196257>
 - <https://github.com/heykeetae/Self-Attention-GAN>
 - <https://xiaosean.github.io/deep%20learning/computer%20vision/2018-06-15-SAGAN/>
- Diffusion models :
 - <https://github.com/tcapelle/Diffusion-Models-pytorch?organization=tcapelle&organization=tcapelle>
- DANN :
 - <https://github.com/fungtion/DANN>
 - https://github.com/NaJaeMin92/pytorch_DANN